

Interfacing C with OCaml

Niki Yoshiuchi

November 30, 2011

Overview

- ▶ Why?
- ▶ A Quick Example
 - ▶ C Code
 - ▶ OCaml Code
 - ▶ Compiling
- ▶ Function signatures
 - ▶ Arity > 5
- ▶ value type
 - ▶ Integral types
 - ▶ Blocks
 - ▶ Structured data
- ▶ Linking
 - ▶ Static
 - ▶ Dynamic
- ▶ What's left?

Why?

- ▶ Bindings to C libraries
- ▶ Multi-core threading
 - ▶ If the C code doesn't require the OCaml runtime we can release the lock
- ▶ Breaking the type system
 - ▶ If you are going to do this, it's probably better to use the Obj module

A Quick Example: C Code

```
1 #include <caml/memory.h>
2 #include <caml/mlvalues.h>
3
4 CAMLprim value caml_add(value a, value b)
5 {
6     CAMLparam2(a, b);
7     int c;
8     c = Int_val(a) + Int_val(b);
9     CAMLreturn(Val_int(c));
10 }
```

A Quick Example: OCaml Code

```
1 external add : int -> int -> int = "caml_add"  
2  
3 let _ =  
4     let three = add 1 2 in  
5     print_int three;  
6     print_newline ();
```

A Quick Example: Compiling

```
1 gcc -I /usr/local/lib/ocaml -c stubs.c
2 ocamlc -c main.ml
3 ar rcs libstubs.a stubs.o
4 ocamlc -o main.byte -custom libstubs.a main.cmo
```

Function Signatures

The OCaml code is simple - declare your function using the *external* keyword, specify its type and set it equal to the name of the C implementation.

► OCaml:

```
1      external name : type = "C-function-name"
```

► C:

```
1      CAMLprim value name(value arg1, value arg2)
```

Arity > 5

For functions with Arity $n > 5$, we must implement two C functions - one with n arguments and one with an array. The former is used for native compilation, the latter for byte code compilation.

► OCaml:

```
1      external name : type = "byte-code-name"  
2      "native-code-name"
```

► C (native):

```
1      CAMLprim value native_code(value arg1, value arg2,  
2      value arg3, value arg4, value arg5, value arg6)
```

► C (byte):

```
1      CAMLprim value byte_code(value *argv, int argc) {  
2          native_code(argv[0], argv[1], argv[2],  
3          argv[3], argv[4], argv[5]);  
4      }
```


value type

All variables passed between OCaml and C are of type *value*. A *value* can be:

- ▶ A integral type (unboxed)
- ▶ A pointer to a heap allocated object
- ▶ An pointer to a object allocated outside of the heap

There are two macros to help you determine what type is encapsulated in a *value*

- ▶ `Is_long(v)` - returns true if `v` is an integral type
- ▶ `Is_block(v)` - returns true if `v` is a pointer to a block

integral types

There are a number of macros for retrieving and storing integral types. They are of the form *Val_type* (think "value of type") and *Type_val* (think "type of value")

- ▶ `Val_long(l)`
- ▶ `Val_int(i)`
- ▶ `Val_bool(b)`
- ▶ `Val_true`
- ▶ `Val_false`
- ▶ `Val_unit`
- ▶ `Long_val(v)`
- ▶ `Int_val(i)`
- ▶ `Bool_val(v)`

Going back to our quick example

Let's take another look at our quick example to see how these macros are used:

```
1 #include <caml/memory.h>
2 #include <caml/mlvalues.h>
3
4 CAMLprim value caml_add(value a, value b)
5 {
6     CAMLparam2(a, b);
7     int c;
8     c = Int_val(a) + Int_val(b);
9     CAMLreturn(Val_int(c));
10 }
```

Blocks

If a *value* is a block, then it has a tag (accessible via the *Tag_val(v)* macro) which will return one of the following values:

Tag	Description
0 to No_scan_tag-1	A structured block. Each field is a <i>value</i> .
Closure_tag	A closure representing a functional value.
String_tag	A character string.
Double_tag	A double precision float.
Abstract_tag	An abstract datatype.
Custom_tag	A custom datatype.

Doubles, strings, Ints

Doubles, strings, Int32s, Int64s and Nativeints are all stored as pointers to heap allocated objects. They are accessible via a set of macros of the form *Type_val*:

- ▶ `Double_val(v)`
- ▶ `String_val(v)`
- ▶ `Int32_val(v)`
- ▶ `Int64_val(v)`
- ▶ `Nativeint_val(v)`

Allocating doubles, strings and Ints

In order to pass a double, string or Int32/Int64/Nativeint from C to OCaml, you must allocate an object on OCaml's heap.

- ▶ `caml_alloc(n, t)` - Allocates a block with tag t and size n
- ▶ `caml_alloc_string(n)` - Allocates a string of length n
- ▶ `caml_copy_string(s)` - Copies the null-terminated string s
- ▶ `caml_copy_double(d)`
- ▶ `caml_copy_int32(i)`, `caml_copy_int64(i)`, `caml_copy_nativeint(i)`

Garbage Collection

If your C code manipulates any heap allocated objects or performs any allocations, you must register it with the garbage collector.

OCaml provides a number of macros:

- ▶ CAMLparam0 to CAMLparam5 - registers 0 to 5 parameters
- ▶ CAMLxparam1 to CAMLxparam5 - for registering additional parameters for functions with arity > 5
- ▶ CAMLlocal1 to CAMLlocal5 - for declaring local *value* variables
- ▶ CAMLlocalN(x, n) - for declaring an array of *values* of size n
- ▶ CAMLreturn0, CAMLreturn(x) and CAMLreturnT(t, x) - these replace the *return* statement

Garbage Collection

- ▶ `CAMLparam#` and `CAMLxparam#` should be the first things called in your function. All your function's parameters should be registered using them. If you have 3 parameters, then use `CAMLparam3`, etc.
- ▶ `CAMLlocal#` and `CAMLlocalN` should be used to declare any local *value* variables used in your function.
- ▶ `CAMLreturn` should replace all the *return* statements in your function. This is required if you used any of the `CAMLparam` or `CAMLlocal` functions.

Double example

```
1 #include <caml/mlvalues.h>
2 #include <caml/memory.h>
3 #include <caml/alloc.h>
4
5 CAMLprim value caml_double_add(value dbl1 ,
6                               value dbl2)
7 {
8     CAMLparam2(dbl1 , dbl2 );
9     CAMLlocal1( ret );
10    double d;
11    d = Double_val(dbl1) + Double_val(dbl2 );
12    ret = caml_copy_double(d);
13    CAMLreturn( ret );
14 }
```

Structured Data

- ▶ Variants
- ▶ Records and tuples
- ▶ Lists
- ▶ Arrays
- ▶ Big Arrays

Variants

- ▶ Constant constructors are represented as unboxed integers. The value is the order in which the constructor appears, starting from 0.
- ▶ Non-constant constructors are represented as blocks. The tag for each block is the order in which the non-constant constructor appears, starting from 0.

Variant Example

```
1 type t =  
2   | Nothing (* Val_int(0) *)  
3   | Int of int (* Block with tag 0 *)  
4   | Double of float (* Block with tag 1 *)  
5   | String of string (* Block with tag 2 *)  
6   | Something (* Val_int(1) *)  
7   | Variant of t * t (* Block with tag 3 *)  
8  
9 external variant : t -> unit = "caml_variant"  
10  
11 let _ =  
12   let v = Variant (Int 5,  
13     Variant (String "Hello",  
14       Nothing)) in  
15   variant v
```

Variant example

```
1 #include <stdio.h>
2 #include <caml/mlvalues.h>
3 #include <caml/memory.h>
4 #include <caml/alloc.h>
5
6 CAMLprim value caml_variant(value t) {
7     CAMLparam1(t);
8     if(Is_block(t)) {
9         switch(Tag_val(t)) {
10             case 0:
11                 printf("Int_of_%d\n", Int_val(Field(t, 0)));
12                 break;
13             case 1:
14                 printf("Double_of_%f\n", Double_val(Field(t, 0)));
15                 break;
16             case 2:
17                 printf("String_of_%s\n", String_val(Field(t, 0)));
18                 break;
19             case 3:
20                 printf("Begin_Variant_of\n");
21                 caml_variant(Field(t, 0));
22                 caml_variant(Field(t, 1));
23                 printf("Variant\n");
24                 break;
25         }
26     }
27     else {
28         if(Int_val(t) == 0)
29             printf("Nothing\n");
30         else if(Int_val(t) == 1)
31             printf("Something\n");
32     }
33     CAMLreturn(Val_unit);
34 }
```

Tuples and Records

Tuples and records are both represented by 0 tagged blocks. Fields are accessed in numerical order using the *Field*(*v*, *n*) macro.

Field(*v*, *n*) returns an lvalue and can be used to both set and get the field. Tuples and records can be allocated using the *caml_alloc_tuple*(*n*) function.

Record example

```
1 type t = {first : string; second : float}
2
3 external create : string -> float -> t =
4     "caml_create"
5
6 let _ =
7     let r = create "Hello" 3.1415 in
8     print_endline r.first;
9     print_float r.second;
10    print_newline ();
```

Record example

```
1 #include <caml/mlvalues.h>
2 #include <caml/memory.h>
3 #include <caml/alloc.h>
4
5 CAMLprim value caml_create(value str, value fl) {
6     CAMLparam2(str, fl);
7     CAMLlocal1(ret);
8
9     ret = caml_alloc_tuple(2);
10    Field(ret, 0) = str;
11    Field(ret, 1) = fl;
12
13    CAMLreturn(ret);
14 }
```


Lists

- ▶ [] is a Val_int of 0
- ▶ Blocks with tag 0
- ▶ Block size of 2
- ▶ Field(v, 0) contains the head
- ▶ Field(v, 1) contains the tail

List example

```
1 external print_list : int list -> unit =  
2     "caml_print_list"  
3  
4 let () = print_list [1;2;3;4;5]
```

List example

```
1 #include <stdio.h>
2 #include <caml/mlvalues.h>
3 #include <caml/memory.h>
4 #include <caml/alloc.h>
5
6 CAMLprim value caml_print_list(value list) {
7     CAMLparam1(list);
8     CAMLlocal2(head, tail);
9
10    if(Is_block(list))
11    {
12        head = Field(list, 0);
13        tail = Field(list, 1);
14        printf("%d_", Int_val(head));
15        CAMLreturn(caml_print_list(tail));
16    }
17
18    printf("\n");
19    CAMLreturn(Val_unit);
20 }
```

Arrays

- ▶ Block with tag 0
- ▶ Use the *Wosize_val(v)* macro to get the size of the array
- ▶ Use *Field(v, n)* to get or set the *n*th element
- ▶ *caml_alloc_tuple(n)* to allocate an array of size *n*
- ▶ Double arrays
 - ▶ Block with tag *Double_array_tag*
 - ▶ Use *Double_field(v, n)* to get the *n*th element
 - ▶ *Store_double_field(v, n, d)* to set the *n*th element
 - ▶ *caml_alloc(n, Double_array_tag)* to allocate an array of size *n*

Array example

```
1 external create_array : int -> int array =  
2     "caml_create_array"  
3 external print_array : int array -> unit =  
4     "caml_print_array"  
5  
6 let _ =  
7     let a = create_array 5 in  
8     print_array a
```

Array example

```
1 #include <stdio.h>
2 #include <caml/mlvalues.h>
3 #include <caml/memory.h>
4 #include <caml/alloc.h>
5
6 CAMLprim value caml_create_array(value size)
7 {
8     CAMLparam1(size);
9     CAMLlocal1(ret);
10    int i;
11    ret = caml_alloc_tuple(Int_val(size));
12    for(i=0; i<Int_val(size); ++i)
13        Field(ret, i) = Val_int(i);
14    CAMLreturn(ret);
15 }
16
17 CAMLprim value caml_print_array(value a)
18 {
19     CAMLparam1(a);
20     int size, i;
21     size = Wosize_val(a);
22     for(i=0; i<size; ++i)
23         printf("%d ", Int_val(Field(a, i)));
24     printf("\n");
25     CAMLreturn(Val_unit);
26 }
```

Double array example

```
1 external create_array : int -> float ->  
2   float array = "caml_create_array"  
3 external print_array : float array -> unit =  
4   "caml_print_array"  
5  
6 let _ =  
7   let a = create_array 5 0.5 in  
8   print_array a
```

Double array example

```
1 #include <stdio.h>
2 #include <caml/mlvalues.h>
3 #include <caml/memory.h>
4 #include <caml/alloc.h>
5
6 CAMLprim value caml_create_array(value size, value inc)
7 {
8     CAMLparam2(size, inc);
9     CAMLlocal1(ret);
10    double d, step;
11    int i;
12    ret = caml_alloc(Int_val(size), Double_array_tag);
13    d = 0.0;
14    step = Double_val(inc);
15    for(i=0; i<Int_val(size); ++i)
16        Store_double_field(ret, i, d + i*step);
17    CAMLreturn(ret);
18 }
19
20 CAMLprim value caml_print_array(value a)
21 {
22     CAMLparam1(a);
23     int size, i;
24     size = Wosize_val(a);
25     for(i=0; i<size; ++i)
26         printf("%f ", Double_field(a, i));
27     printf("\n");
28     CAMLreturn(Val_unit);
29 }
```


Bigarrays

- ▶ Native to C (or Fortran)
- ▶ pack data efficiently
- ▶ Are limited to only certain types, mainly integers, floats, chars and complex numbers
- ▶ in OCaml there are three different Bigarrays:
 - ▶ Array1 - Single dimension arrays
 - ▶ Array2 - Two dimensional arrays
 - ▶ Genarray - Multi-dimensional arrays
- ▶ C doesn't care. You get a `void *` and `num_dims` and flags

What do Bigarrays look like?

This is from OCaml 3.12.1:

```
1 struct caml_ba_array {  
2     void * data;  
3     intnat num_dims;  
4     intnat flags;  
5     struct caml_ba_proxy * proxy;  
6     intnat dim[1] ;  
7 };
```

Bigarray - types

- ▶ CAML_BA_FLOAT32 - Single-precision floats
- ▶ CAML_BA_FLOAT64 - Double-precision floats
- ▶ CAML_BA_SINT8 - Signed 8-bit integers
- ▶ CAML_BA_UINT8 - Unsigned 8-bit integers
- ▶ CAML_BA_SINT16 - Signed 16-bit integers
- ▶ CAML_BA_UINT16 - Unsigned 16-bit integers
- ▶ CAML_BA_INT32 - Signed 32-bit integers
- ▶ CAML_BA_INT64 - Signed 64-bit integers
- ▶ CAML_BA_CAML_INT - Caml-style integers (signed 31 or 63 bits)
- ▶ CAML_BA_NATIVE_INT - Platform-native long integers (32 or 64 bits)
- ▶ CAML_BA_COMPLEX32 - Single-precision complex
- ▶ CAML_BA_COMPLEX64 - Double-precision complex

Bigarray - layouts and flags

OCaml supports both row-major (C style) and column-major (Fortran style) arrays. This is specified by one of two flags:

- ▶ `CAML_BA_C_LAYOUT`
- ▶ `CAML_BA_FORTRAN_LAYOUT`

Bigarrays also have flags specifying how the memory is managed:

- ▶ `CAML_BA_EXTERNAL` - Data is not allocated by Caml
- ▶ `CAML_BA_MANAGED` - Data is allocated by Caml
- ▶ `CAML_BA_MAPPED_FILE` - Data is a memory mapped file

Bigarrays - allocation

There are two functions provided to allocate Bigarrays in C:

- ▶ value `caml_ba_alloc(int flags, int num_dims, void * data, intnat * dim);`
 - ▶ `flags` - The type and layout ORed together
 - ▶ `num_dims` - the number of dimensions
 - ▶ `data` - a pointer to the C array
 - ▶ `dim` - a pointer to an array containing the size of each dimension
- ▶ value `caml_ba_alloc_dims(int flags, int num_dims, void * data, ...);`
 - ▶ Pretty much the same, except it takes a `va_args` list of dimensions

Bigarrays - access

There are two macros - one returns the Bigarray struct, the other the data pointer

- ▶ `Caml_ba_array_val(v)` - returns the struct
- ▶ `Caml_ba_data_val(v)` - returns the data pointer

Bigarray example

```
1 open Bigarray
2
3 external create_ba : int -> int ->
4     (int, int64_elt, c_layout) Array2.t
5     = "caml_create_ba"
6
7 external print_ba :
8     (int, int64_elt, c_layout) Array2.t -> unit
9     = "caml_print_ba"
10
11 let _ =
12     let ba = create_ba 5 5 in
13     print_ba ba
```

Bigarray example

```
1 #include <stdio.h>
2 #include <caml/mlvalues.h>
3 #include <caml/memory.h>
4 #include <caml/alloc.h>
5 #include <caml/bigarray.h>
6
7 CAMLprim value caml_create_ba(value w, value h)
8 {
9     CAMLparam2(w, h);
10    CAMLlocal1(ret);
11    int i, width, height;
12    long *data;
13
14    width = Int_val(w);
15    height = Int_val(h);
16    /* oh noes! leaking memory! */
17    data = malloc(sizeof(long)*width*height);
18    for(i=0; i<width*height; ++i)
19        data[i] = i;
20    ret = caml_ba_alloc_dims(CAML_BA_INT64 | CAML_BA_C_LAYOUT,
21                             2, data, width, height);
22    CAMLreturn(ret);
23 }
24
25 CAMLprim value caml_print_ba(value ba)
26 {
27     CAMLparam1(ba);
28     int width, height;
29     int x, y;
30     long *data;
31     width = Caml_ba_array_val(ba)->dim[0];
32     height = Caml_ba_array_val(ba)->dim[1];
33     data = Caml_ba_data_val(ba);
34     for(y=0; y<height; ++y) {
35         for(x=0; x<width; ++x) {
36             printf("%d_", data[y*width + x]);
37         }
38         printf("\n");
39     }
40    CAMLreturn(Val_unit);
41 }
```


Linking

Use Oasis.

Static Linking

To build the C library:

```
1 gcc -I /usr/local/lib/ocaml -c stubs.c
2 ar rcs libstubs.a stubs.o
```

To link the OCaml code:

```
1 ocamlc -o main.byte -custom libstubs.a main.cmo
2 ocamlopt -o main.native libstubs.a main.cmx
```

Dynamic Linking

To build the C library:

```
1 gcc -I /usr/local/lib/ocaml -c -fPIC  
2     stubs.c -o stubs.o  
3 gcc -shared -Wl,-soname,libstubs.so -o libstubs.so
```

To link the OCaml code:

```
1 ocamlc -o main.byte libstubs.so main.cmo  
2 ocamlopt -o main.native libstubs.so main.cmx
```

What's left?"

- ▶ Callbacks
- ▶ Exceptions
- ▶ Custom data
- ▶ Non-blocking calls
- ▶ C threads

Fin.

Thanks for coming! The slides and code are available on github:
<https://github.com/aplusbi/OCaml-FFI-presentation>